

## DS4300: FOUNDATIONS

### Searching

- Searching is the most common operation in a database system.
- **SQL SELECT Statement:** The most versatile and complex SQL command used for querying data.
- **Baseline Efficiency: Linear Search**
  - Starts at the beginning of a list.
  - Proceeds element by element until:
    - The desired item is found.
    - The end of the list is reached without finding the item.
- **Record** = A collection of values for attributes of a single entity instance; a row of a table
- **Collection** = a set of records of the same entity type; a table
- Trivially, stored in some sequential order like a list
- **Search Key** = A value for an attribute from the entity type
  - Could be  $\geq 1$  attribute

### List of Records

- If each record takes up  $x$  bytes of memory, then for  $n$  records, we need  $n*x$  bytes of memory.
- *Contiguously Allocated List*
  - All  $n*x$  bytes are allocated as a single “chunk” of memory
- *Linked List*
  - Each record needs  $x$  bytes + additional space for 1 or 2 memory addresses
  - Individual records are linked together in a type of chain using memory addresses

### Contiguous vs. Linked Allocation

- *Contiguous*
  - Data is stored in **one continuous block** of memory.
  - Example: **Arrays** in programming.
  - **Pros:**
    - Fast access for random access
  - **Cons:**
    - Slow for random insertions except at the end of the array
- *Linked*
  - Data is stored in **separate, non-adjacent memory locations**, with pointers connecting them.
  - Example: **Linked Lists** in programming.
  - **Pros:**
    - Faster for random insertion
  - **Cons:**
    - Slower access

## Binary Search

- **Input** = Array of values in sorted order, target value
- **Output** = The location (index) of where the target is located or some value indicating the target was not found
- **Sample code:**

```
def binary_search(arr, target)
    left, right = 0, len(arr) - 1
    while left <= right:
        mid = (left + right) // 2
        if arr[mid] == target:
            return mid
        elif arr[mid] < target:
            left = mid + 1
        else:
            right = mid - 1
    return -1
```

## Time Complexity

- **Linear Search**
  - **Best case:** target is found at the first element; only 1 comparison
  - **Worst case:** Target is not in the array; n comparisons
  - Therefore, in the worst case, linear search is  $O(n)$  time complexity
- **Binary Search**
  - Best case: Target is found at mid; 1 comparison (inside the loop)
  - Worst case: Target is not in the array;  $\log_2 n$  comparisons.
  - Therefore, in the worst case, binary search is  $O(\log_2 n)$  time complexity.

## Back to Database Searching

- Assume data is stored on disk by column *id*'s value
- Searching for a specific *id* = fast
- A linear scan is the only option when searching for a specific *specialVal*.
- Data can't be stored on a disk sorted by both *id* and *specialVal* simultaneously.
- Sorting by both would require duplicating data, which is space-inefficient.
- **Solution:** Use an external data structure to enable faster searching for *specialVal* instead of relying on a linear scan
- **Options:**
  - An array of tuples (*specialVal*, *rowNumber*) sorted by *specialVal*

- We could use Binary Search to quickly locate a particular specialVal and find its corresponding row in the table
- But, every insert into the table would be like inserting into a sorted array - slow...
- A linked list of tuples (specialVal, rowNum) sorted by specialVal
  - searching for a specialVal would be slow - linear scan required
  - But inserting it into the table would theoretically be quick to also add to the list.
- Something with fast insert and fast search:
  - **Binary Search Tree** = a binary tree where every node in the left subtree is less than its parent and every node in the right subtree is greater than its parent.